

BHUSHAN JAIN

Azure Data Engineer

Phone: 7066858602 | Email: bhushanjain311997@gmail.com | Pune, India

PROFESSIONAL SUMMARY

Azure Data Engineer with 5.9+ years of experience designing, building, and operating cloud data platforms on Microsoft Azure for telecom-scale workloads. Skilled in building end-to-end batch and streaming data pipelines using Azure Data Factory, Apache Airflow, and PySpark on Azure Databricks, with strong SQL and data modeling expertise (star/snowflake schemas, ERDs, fact/dimension design). Experienced integrating Oracle and PostgreSQL sources into Azure Data Lake Storage (ADLS Gen2), implementing Medallion (Bronze-Silver-Gold) architectures with Unity Catalog, and delivering data to Azure SQL and downstream analytics platforms. Hands-on with CI/CD via Azure DevOps and Git, and monitoring/alerting via Azure Monitor, Log Analytics, and Application Insights in an SLA-driven operating model. Working knowledge of core Java. Comfortable working cross-functionally with architects and analytics teams to translate business requirements into scalable data solutions.

WORK EXPERIENCE

Azure Data Engineer — Amdocs

Jun 2022 – Present

- Deliver Azure data engineering for US Cellular's M2M billing platform (APRM): pipeline design, optimization, defect management, and cloud data architecture in an SLA-driven operating model.
- Designed and deployed end-to-end ADF pipelines ingesting multi-source data into ADLS Gen2, reducing latency and manual intervention.
- Built and optimized PySpark transformation frameworks in Azure Databricks to process 2TB+ of daily high-velocity data for large-scale retail, sales, and customer datasets, improving processing efficiency by 20%.
- Integrated Oracle DB and PostgreSQL sources into Azure using ODI and external tables for lake ingestion; built and tuned Azure SQL data structures, improving query performance through indexing and schema design.
- Automated cross-functional workflows using Azure Data Factory and Apache Airflow, reducing manual intervention by 30%.
- Optimized PySpark transformation logic and partitioning strategies, reducing cloud compute costs by 15% while handling large-scale joins across billing and customer tables.
- Implemented monitoring and alerting with Azure Monitor, Log Analytics, and Application Insights, reducing pipeline downtime by 25% and accelerating reporting from 24 hours to near real-time.
- Architected end-to-end Medallion pipelines (Bronze-Silver-Gold) using Unity Catalog to centralize metadata and enforce fine-grained access control across workspaces.
- Collaborated with BI/analytics teams for downstream reporting and used Git and Azure DevOps for controlled deployments across environments.

Associate Engineer — TCL

May 2020 – May 2022

- Orchestrated scalable batch and streaming pipelines using Azure Data Factory (ADF) to ingest high-volume transactional data into ADLS Gen2.
- Developed real-time Spark Streaming applications to ingest and process Kafka data, with window-based aggregations to identify real-time trends and patterns.
- Developed optimized PySpark and Python scripts in Databricks to clean, transform, and aggregate complex structured and unstructured datasets.
- Formulated high-performance SQL queries to implement incremental loading and Delta Lake architecture, improving pipeline efficiency.
- Designed and implemented star and snowflake schemas for CRM and OMS applications, integrating modules like Customer Profile, Billing, Payments, and Transactions.
- Built fact and dimension tables with well-defined primary and foreign key relationships, ensuring referential integrity to support complex analytical queries; developed ER diagrams (ERDs) to improve cross-team collaboration.

- Automated CI/CD release pipelines using Azure DevOps (Git-based), enabling seamless code deployment across development, staging, and production environments.
- Monitored pipeline health and diagnosed runtime failures using Azure Monitor and Log Analytics; configured automated alerting for proactive issue resolution.

SKILLS

Cloud & Orchestration: Azure Data Factory (ADF), Azure Databricks, Azure Data Lake Storage (ADLS Gen2), Azure Synapse Analytics, Apache Airflow

Data Engineering: PySpark, Spark SQL, Apache Spark, Kafka (Spark Streaming), ETL/ELT Pipelines, Delta Lake, Medallion Architecture, Unity Catalog

Programming: Python (Pandas, data processing), SQL, Core Java (basics)

Data Modeling: Star & Snowflake Schemas, Data Modeling, Normalization/Denormalization, ERDs

Databases: Azure SQL, Oracle, PostgreSQL

DevOps & Monitoring: Azure DevOps, Git, CI/CD, Azure Monitor, Log Analytics, Application Insights

Tools: Jenkins, GitHub, YARN

EDUCATION

- M.Tech, Computer Engineering — VIIT, Pune | 2018 – 2020
- B.E., Computer Science & Engineering — D.N. Patel College of Engineering, Shahada (NMU) | 2015 – 2018
- Diploma — Government Polytechnic, Nandurbar (MSBTE) | 2012 – 2015

CERTIFICATIONS

- AWS Certified Cloud Practitioner — Amazon Web Services (Valid Apr 2025 – May 2029)
- Oracle Cloud Infrastructure 2024 Generative AI Certified Professional
- Microsoft Azure AI Fundamentals (AI-900)

AWARDS & RECOGNITION

- Employee of the Month — Amdocs: Outstanding production support, bug analysis, and incident resolution.
- Client Appreciation Award — USCC (US Cellular) for rapid resolution of high-priority production defects.
- Townhall Recognition — Senior Management for operational efficiency and quality delivery.